



Benchmarking

CPU/MEMORY/DISK AND NETWORK

Chetan Gupta | CS553 | 30/03/2018

1: INTRODUCTION

Benchmarking of the components of the processor, memory, network and hard disk is complete. Several experiments are performed on different parameters. Three tests for each experiment are completed, and the mean and average deviation of the results is mentioned. Each experiment takes several seconds to reduce the effect of the initial load or any other system noise. The general values are used on the graphs. This document illustrates the design strategy for CPU / memory / disk/ network and has a corresponding section under each of them, showing diagrams / trends.

1.1 GENERAL POINTS

We have created a single interface for calling all test checks based on command line parameters. Benchmarking code for memory, network and disk is written in C language and CPU is written in C programming.

1.2 TESTING PLATFORM

I have use used my CPU and Hyperian cluster for testing my codes, details are as follows:-

2: CPU Benchmarking

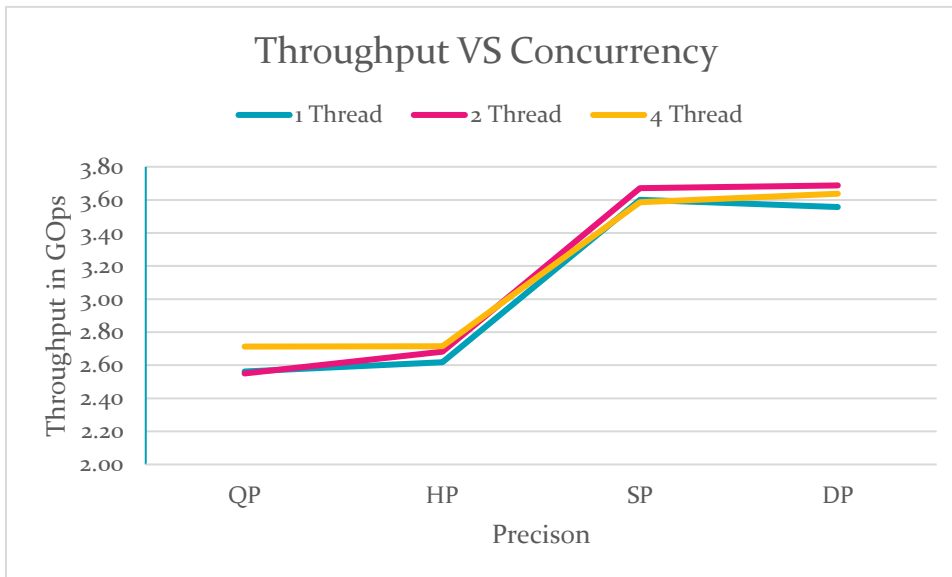
For CPU benchmarking we have calculated both IOPS and FLOPS with the help of C programming language. Following steps are followed for the calculation:

- An input is taken from the user about the number of threads which is to be assigned for the benchmarking.
- The user can choose 1, 2, or 4 threads from options for calculating IOPS/FLOPS(GOPS) of CPU. The threads are parallelized to perform the operations.
- For calculating the IOPS/FLOPS all the threads do the same calculation for $(2^{31}-1)$ times and the total time taken for performing these operations are recorded.
- Each thread are doing same set of calculations one Trillion times. Total time taken for these operations is recorded, and this is used for calculating the IOPS / FLOPS(GOPS).
- The sums are calculated, and these are used for plotting the graph using GNU plot.

		Measured Ops/Sec	Measured Ops/Sec	Ops/Sec (GigaOPS)	Efficiency (%)	Efficiency (%)
--	--	---------------------	---------------------	----------------------	----------------	-------------------

		(GigaOPS)	(GigaOPS)	(Hyperion)		
QP	1	2.562159	N/A	588.8	0.44	N/A
QP	2	2.5503425	N/A	588.8	0.43	N/A
QP	4	2.713115	N/A	588.8	0.46	N/A
HP	1	2.617619	N/A	294.4	0.89	N/A
HP	2	2.681138667	N/A	294.4	0.91	N/A
HP	4	2.7148875	N/A	294.4	0.92	N/A
SP	1	3.6003485	N/A	147.2	2.45	N/A
SP	2	3.67169	N/A	147.2	2.49	N/A
SP	4	3.5848745	N/A	147.2	2.44	N/A
DP	1	3.557330667	35.25	73.6	4.83	10.09
DP	2	3.6872475	63.8	73.6	5.01	5.78
DP	4	3.637126667	70.12	73.6	4.94	5.19

2.1.1 PERFORMANCE CHART: GIOPS/GFLOPS VS THREADS



Observations:

- GOPS values increases as tnumber of threads increases than becomes constant .

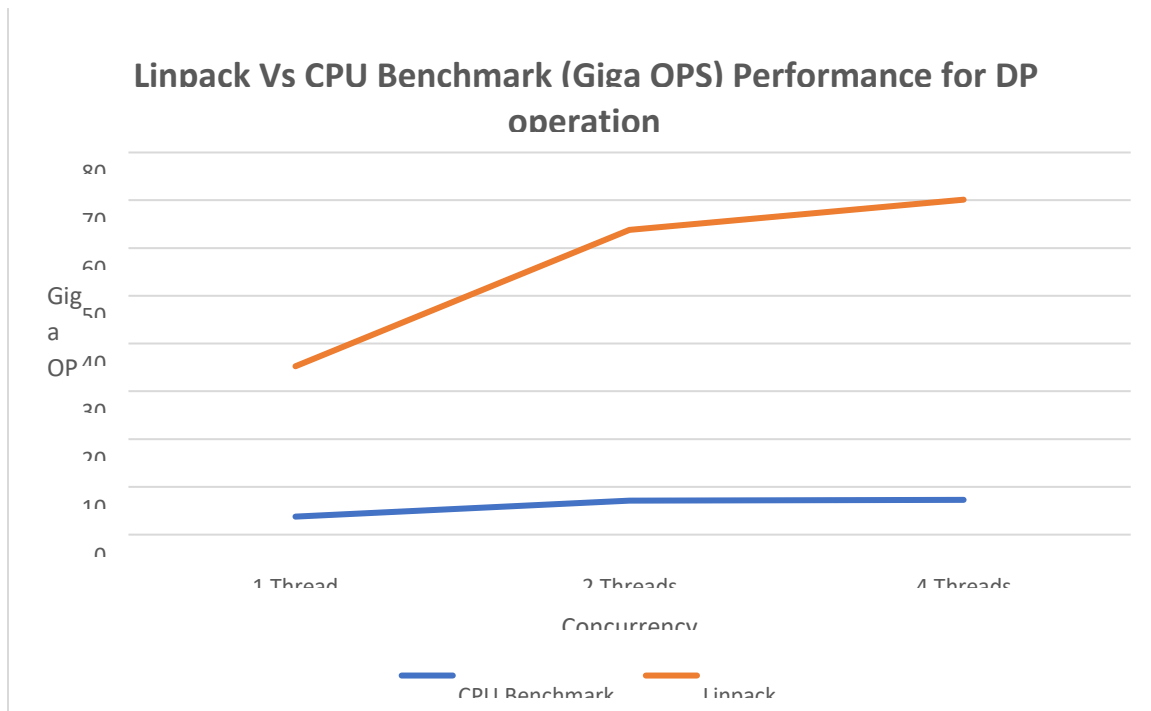
We used 1, 2, 4 threads for this experiment and 4 readings of GOPS is recorded and then plotted. To calculate the height performance we ran $2^{31}-1$ arithmetic operations which include addition, multiplication and

subtraction. It also depicts that the increase in the number of threads improve GigaOPS to double the size of threads. It also shows that for the threads count 2 and 4, the GigaOPS tend to approximately similar; the reason being is that given number of threads similar to the system available threads (or higher than system available threads) provides or thus gives low GigaOPS sometimes due to thread switching.

Conclusion-

It can be easily concluded from above bar graphs that the GigaOPS for single thread are almost half than that of two threads for each QP, HP, SP, and DP operations.

Below line graph represent comparison of Linpack and CPU benchmark DP floating point operation.



Conclusion-

From above graph, we can see that **Linpack** benchmark gives highest performance for for DP floating point operations. The result of **Linpack** benchmark also follows the same pattern in growth of GigaOPS for the given number of threads similar to that of CPU benchmark measured on Hyperion. The reason behind high performance of Linpack benchmark is because of high performance CPU utilization and optimum implementation of parallelism.

Processor Performance Metrics:

GigaOPS = (No of Operations / Time) / 10^9

Theoretical OPS = $2.3 * 2 * \text{instructions per second}$

Theoretical OPS QP= 588.80 GigaOPS

Theoretical OPS HP= 294.40 GigaOPS

Theoretical OPS SP= 147.20 GigaOPS

Theoretical OPS DP= **73.60 GigaOPS**

Maximum Efficiency is achieved by QP operations when the concurrency is 2 threads and efficiency is **18.010923 %**.

Performance Comparison between CPU Benchmark, Linpack Benchmark and Theoretical peak performance

GigaOPS measured by CPU Benchmark for all operations follow the same growth rate as the number of threads increase which is alike to the Linpack performance in growth rate which gets double when thread size increases from 1 to 2 and remains approximately similar when threads size increases from 2 to 4. However, theoretical throughput is computed based on number of instructions of each operation multiplied by number of cores into system and CPU speed.

3. Memory Benchmarking:

For Memory benchmarking we measure the speed of memory with the help of copy operations which are read and write with different block sizes such as 8b, 8KB, 8MB and 80MB. Following steps are followed for the calculation:

- a. An input is taken from the user about the memory block size (1 Byte) or (1KB) or (1 MB) or (10MB)), and number of threads are also asked by the user (1 or 2 or 4) to calculate the memory speed and the space is allocated by read and write.
- b. Access type such as Sequential or Random are taken as input from the user for memory benchmark.
- c. Block by block copy is performed for the sequential access.
- d. Copy is performed based on randomly derived blocks size number for random access.
- e. Memory blocks are distributed between threads for multiple thread operations and each thread works on one half of the fully allocated memory space.
- f. For the evaluation of the Memory benchmark, I used GBps for the measuring the throughput of the memory.
- g. I ran the Memory benchmark for concurrency of 1,2,4 threads with Read-Write Sequential and Read-Write Random access patterns with workload of 100GB data with block sizes of 1KB, 1MB and 10MB.

h. I also used the rand () function from *math.h* library to generate the random numbers needed for RWR access pattern.

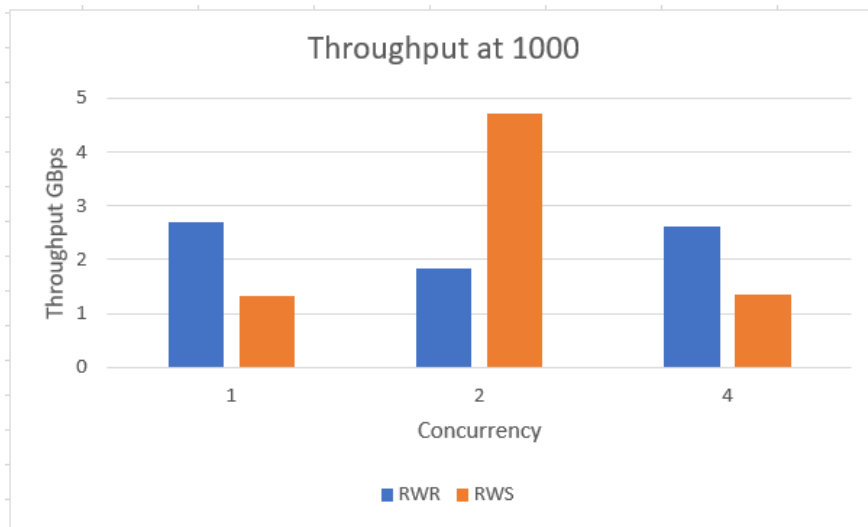
i. For comparison, I ran the *pmbw* benchmark to measure memory subsystem performance presented in tabular format below.

Below table 2 shows the Memory Throughput measured on Hyperion Cluster.

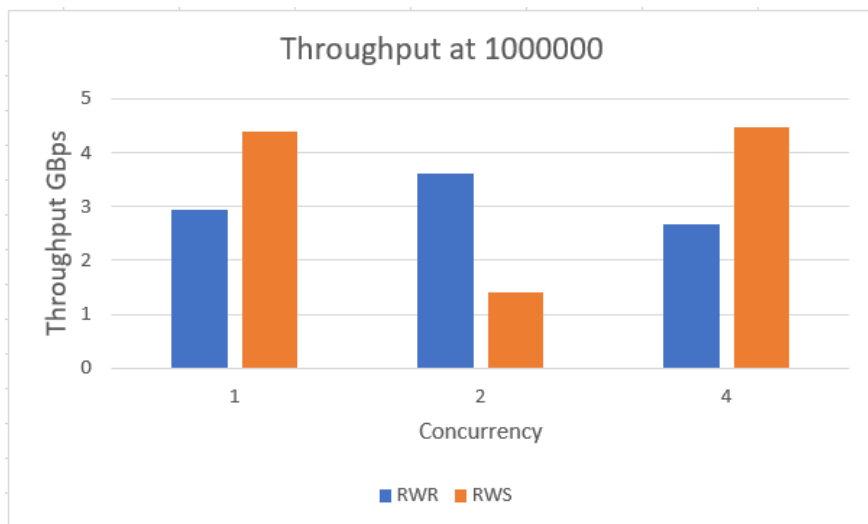
3.1 Performance Chart and details:

Workload	ConCurrency	Block	MyRAMBench	Pmbw	Theoretical	MyRAMBench	pmbw
		Size	Measured	Measured	Throughput	Efficiency (%)	Efficiency
			Throughput	Throughput	(GB/sec)		(%)
			(GB/sec)	(GB/sec)			
RWS	1	1000	2.216058	16.5	68.256	3.246686006	24.173699
RWS	1	1000000	1.416238667	20.7	68.256	2.074892561	30.3270042
RWS	1	10000000	2.387083333	19.9	68.256	3.497250547	29.1549461
RWS	2	1000	1.208943	24.9	68.256	1.771189346	36.4803094
RWS	2	1000000	2.129539667	15	68.256	3.11993036	21.97609
RWS	2	10000000	1.351341667	36	68.256	1.97981374	52.742616
RWS	4	1000	4.797787	22.6	68.256	7.029106599	33.1106423
RWS	4	1000000	1.337346667	34.6	68.256	1.959310048	50.6915143
RWS	4	10000000	4.783803333	38.7	68.256	7.008619511	56.6983122
RWR	1	1000	1.815215667	4.2	68.256	2.659422859	6.1533052
RWR	1	1000000	6.453729	0.61	68.256	9.455181962	0.89369433
RWR	1	10000000	2.890362667	0.49	68.256	4.234591342	0.71788561
RWR	2	1000	4.818662667	8.9	68.256	7.059690967	13.0391467
RWR	2	1000000	2.886603	1.24	68.256	4.229083158	1.81669011
RWR	2	10000000	5.943713667	0.85	68.256	8.707972437	1.24531177
RWR	4	1000	1.777489333	19.4	68.256	2.604151039	28.4224098
RWR	4	1000000	6.099959	3.01	68.256	8.936883204	4.40986873
RWR	4	10000000	2.995375667	0.91	68.256	4.388443018	1.33321613

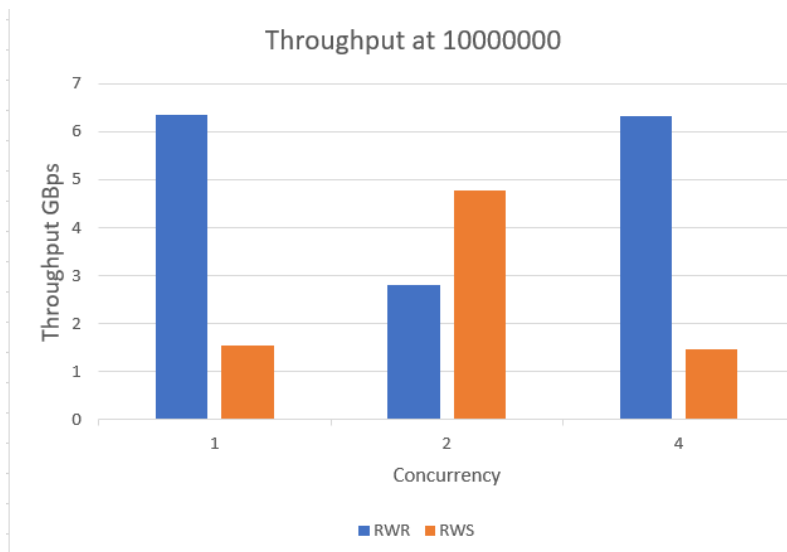
Throughput for Sequential and Random operation in Log plot: Throughput (Y axis) in Mbps with different Operation (X axis) at 1000Bytes of Block size



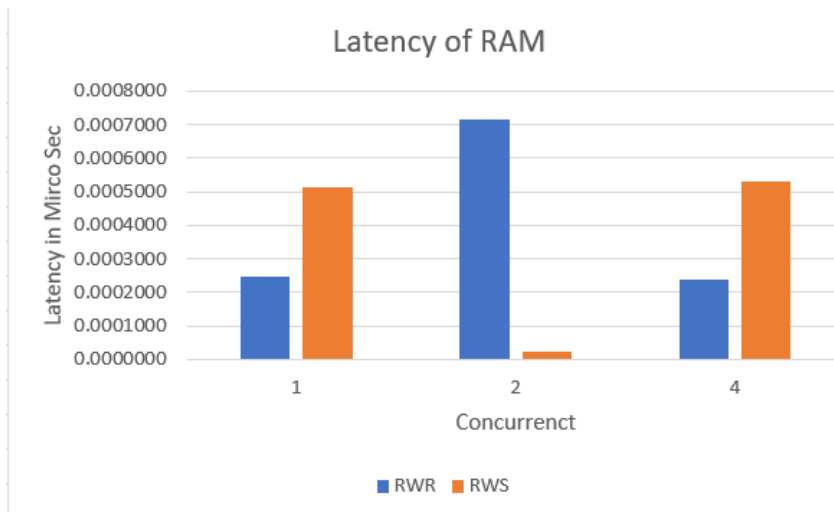
Throughput for Sequential and Random operation in Log plot: Throughput (Y axis) in Mbps with different Operation (X axis) at 100000Bytes of Block size



Throughput for Sequential and Random operation in Log plot: Throughput (Y axis) in Mbps with different Operation (X axis) at 1000000Bytes of Block size



Throughput for Sequential and Random operation in Log plot: Latency (Y axis) in Microseconds with different concurrency (X axis) at 1Bytes of Block size



Observations:

- Read /write speed in sequential access is much more compared to random access. This is because sequential access will have more cache hits as data is accessed sequentially
- Both random and sequential access performed best with block size of 8KB.

Conclusion-

Above tables gave us brief about that the throughput and latency for RWS and RWR access patterns. For RWS access pattern, throughput gets doubled for given block sizes as number of threads increase from 1 to 2, remains approximately similar when increased from 2 to 4 threads.

Additionally, RWS gives us less latency as compared to the *pmbw* measured and theoretical latency. Above tables gave us brief about that the throughput and latency for RWS and RWR access patterns. For RWR access pattern, throughput gets doubled for block sizes of 1MB and 10 MB as number of threads increase from 1 to 2, remains approximately similar when increased from 2 to 4 threads. In addition to this, RWR gives us high latency as compared to theoretical latency but less than *pmbw* measured, the reason being decrease in latency efficiency is that it involves randomly reading and writing data to memory which increases execution time. The reason behind improvement in the throughput is the block size which if increased causes decrease in the memory read and write operations. For block size of 1KB, RWR seems to be not changing considering its throughput for given number of threads.

Conclusion:

From the above comparison graphs, we can say growth in number of threads decreases latency for sequential read-write access pattern because concurrent multiple threads behavior. But growth in number of threads increases latency for random read-write access pattern due to selection of random location for reading data and writing it to memory which will increase the overall latency.

Memory Performance Metrics:

Workload = 1GB

Throughput = $((100 * \text{Workload}) / \text{Time}) / \text{Workload GBps}$

Theoretical Throughput = $(2133 * 2 * 64 * 2) / 8) / 1000 \text{ GBps}$

Throughput Efficiency = $(\text{Throughput} / \text{Theoretical Throughput}) * 100$

Latency = $(\text{execution time} * 1000000) / 100000000 \text{ microseconds}$

Theoretical Latency = 0.014 microseconds

Latency Efficiency = $((\text{Theoretical Latency} - \text{Latency}) / \text{Theoretical Latency}) * 100$

Maximum Throughput Efficiency is achieved by RWR access pattern when the concurrency is 2 threads and efficiency is **21.767445 %**.

Maximum Latency Efficiency is achieved by RWS access pattern when the concurrency is 4 threads and efficiency is **79.19 %**.

Performance Comparison between Memory Benchmark, pmbw Benchmark and Theoretical performance

Throughput measured by Memory Benchmark for RWS access pattern increases as the number of threads increase which gets double when thread size increases from 1 to 2 and

remains approximately similar when threads size increases from 2 to 4. RWR also follows the same pattern that of the RWS except for block size 1KB. **pmbw** benchmark seems to have variations in the throughput for RWS access pattern but for RWR access pattern generates higher throughput for 1 thread and significantly decreases as the number of threads increases. For large block sizes, we get higher throughput due to less number of computations/operations to perform.

Growth rates in threads affect the performance of memory that decreases latency for RWS access pattern but increases for RWR access pattern.

4. Disk Benchmarking:

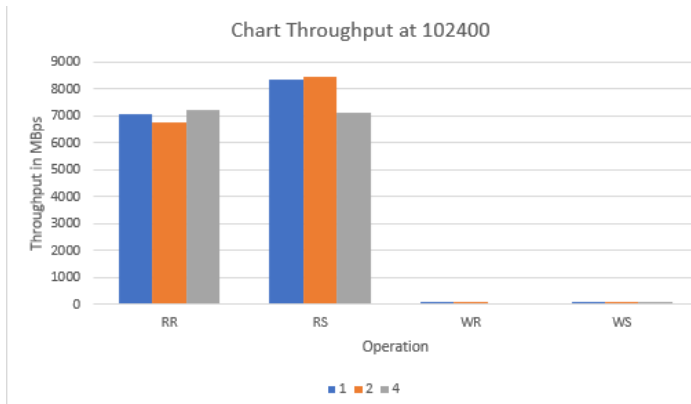
For Disk benchmarking, we analyze multiple read and write operations on disk in both sequential and random fashion. User also has an option to select different numbers of thread to analyze the change in the performance with different threads. Different block sizes of data is written and also read. Following steps are followed for the calculation:

- a. For operations such as read/write, block size, sequential/random access and number of threads, the user will give the input for the operations.
- b. For threads to do simultaneous read/write, create the required number of threads to parallelize the operations.
- c. 10GB space should be provided for the operations to be performed.
- d. Same string of data is read/written in a loop on all over the space of 10GB for sequential read/write.
- e. As compared to sequential experiments the random access is slow, the number of read/write operations are less. The disk space is accessed in a random fashion which is 10GB.
- f. When multiple thread are involved, each thread write to/ read from its own file of 10GB size to work on.
- g. I ran the disk benchmark for concurrency of 1,2,4,8,16,32,64,128 threads with Random Read and Random Write access pattern while concurrency of 1,2,4 threads for Read Sequential and Write Sequential access patterns with workload of 10GB data with block sizes of 1KB, 1MB, 10MB and 100MB.
- h. I also used the `rand ()` function from `math.h` library to generate the random numbers needed for RR and WR access patterns.
- i. For comparison, I ran the **iozone** benchmark to measure disk performance presented in tabular format below.

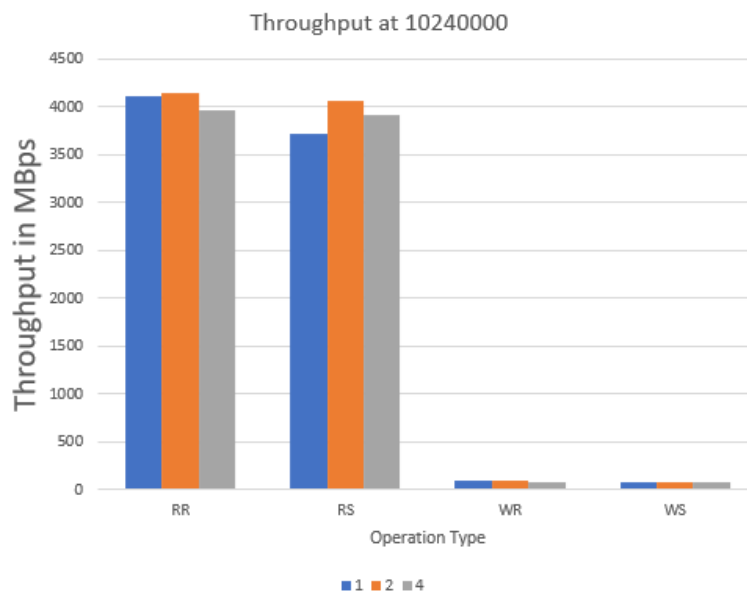
4.1 Performance measurement:

Work-Load	Concurrency	Block Size	MyDiskBench Measured	IOZone Measured	Theoretical Throughput (MB/sec)	MyDiskBench Efficiency (%)	IOZone Efficiency (%)
			Throughput (MB/sec)	Throughput (MB/sec)			
RS	1	1024000	83.3673999	250	372	22.41059137	67.204301
RS	1	10240000	57.60887207	250	372	15.48625593	67.204301
RS	1	102400000	59.64	250	372	16.03225806	67.204301
RS	2	1024000	64.25495484	266	744	8.636418661	35.752688
RS	2	10240000	40.62854859	276	744	5.460826423	37.096774
RS	2	102400000	41.72107056	296	744	5.607670774	39.784946
RS	4	1024000	55.24890625	276	1488	3.71296413	18.548387
RS	4	10240000	40.92319458	290	1488	2.75021469	19.489247
RS	4	102400000	36.47611328	300	1488	2.451351699	20.16129
WS	1	1024000	0.77867577	255	172	0.452718471	148.25581
WS	1	10240000	0.80457352	250	172	0.467775302	145.34883
WS	1	102400000	0.815240635	230	172	0.473977113	133.72093
WS	2	1024000	0.779047695	299	344	0.226467353	86.918604
WS	2	10240000	0.78782349	315	344	0.229018456	91.569767
WS	2	102400000	0.792292095	360	344	0.230317469	104.65116
WS	4	1024000	0.812079695	350	688	0.118034839	50.872093
WS	4	10240000	0.776237605	370	688	0.112825233	53.779069
WS	4	102400000	0.81950409	360	688	0.119113967	52.325581
RR	1	1024000	70.82223145	135	540	13.11522805	25
RR	1	10240000	58.89396729	148	540	10.90629024	27.407407
RR	1	102400000	51.6678003	146	540	9.568111167	27.037037
RR	2	1024000	49.92111694	159	1080	4.622325643	14.722222
RR	2	10240000	41.38459595	169	1080	3.831907032	15.648148
RR	2	102400000	39.71565064	170	1080	3.677375059	15.74074
RR	4	1024000	50.24557374	160	2160	2.326183969	7.4074074
RR	4	10240000	40.79599243	162	2160	1.888703353	7.5
RR	4	102400000	37.92998291	165	2160	1.756017727	7.6388888
WR	1	1024000	0.77867577	240	410	0.18992092	58.536585
WR	1	10240000	0.84928627	242	410	0.207142993	59.02439
WR	1	102400000	0.84145653	248	410	0.2052333	60.487804
WR	2	1024000	0.772429925	280	820	0.094198771	34.146341
WR	2	10240000	0.83750626	283	820	0.10213491	34.512195
WR	2	102400000	0.86670681	305	820	0.105695952	37.195121
WR	4	1024000	0.475235985	201	1640	0.028977804	12.256097
WR	4	10240000	0.84975918	210	1640	0.051814584	12.804878
WR	4	102400000	0.927718085	211	1640	0.056568176	12.865853

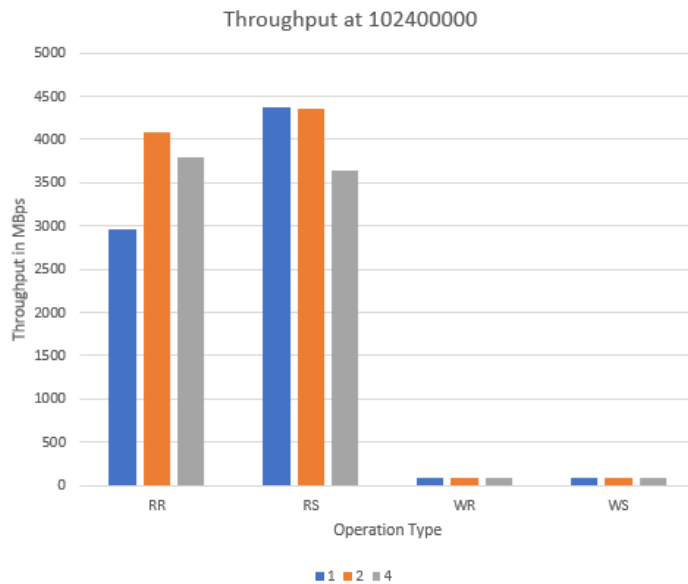
Throughput in Mbps (Y axisà log plot) for all operations at 102400 blocksize



Throughput in Mbps (Y axisà log plot) for all operations at 10240000 blocksize



Throughput in Mbps (Y axisà log plot) for all operations at 102400000 blocksize



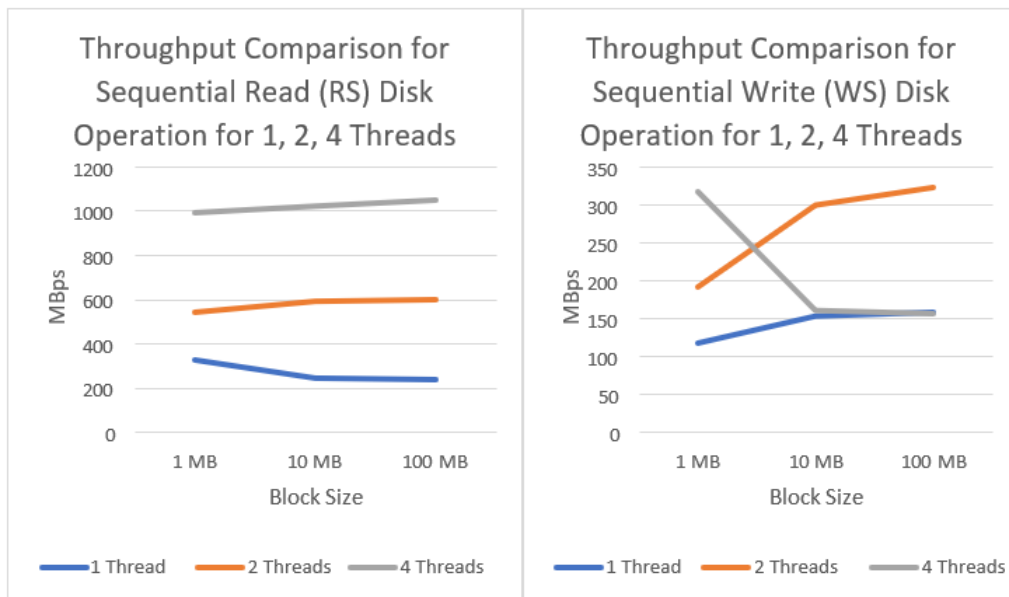
Conclusion-

The table above along with the bar graphs show us that as the number of thread increases up to number of available threads in the system, the throughput increases. We can also see that as the block size increases the throughput increases, the reason being that increase in block size decreases number of disk read or write operations

Below line graphs represent throughput (Mbps) comparison for RS, WS, RR, WR against the concurrency of 1, 2 and 4 threads.

Conclusion –

As can be seen from above line charts that for RS and RR operations, throughput increases with the increase in threads. For WS and WR operations, however throughput varies.



Conclusion –

As can be seen from above line charts that for RS and RR operations, throughput increases with the increase in threads. For WS and WR operations, however throughput varies.

Workload	Concurrency	Block Size	MyDiskBenchmark Measured Latency(ms)	IOZONE Measured Latency(ms)	Theoretical Latency(ms)	MyDiskBenchmark Efficiency	IOZONE Efficiency
RR	1	1KB	0.0329675	1.3	0.5	6.5935	-160
RR	2	1KB	0.0315535	1.3	0.5	6.3107	-160
RR	4	1KB	0.0319245	2	0.5	6.3849	-300
RR	8	1KB	0.034287	1.2	0.5	6.8574	-140
RR	16	1KB	0.030365	0.6	0.5	6.073	-20
RR	32	1KB	0.033113	1.1	0.5	6.6226	-120
RR	64	1KB	0.0408095	1.1	0.5	8.1619	-120
RR	128	1KB	0.0335895	0.56	0.5	6.7179	-12
WR	1	1KB	0.1087795	1.9	0.5	21.7559	-280
WR	2	1KB	0.1116235	1.9	0.5	22.3247	-280
WR	4	1KB	0.1113145	3.1	0.5	22.2629	-520
WR	8	1KB	0.120874	0.75	0.5	24.1748	-50
WR	16	1KB	0.132994	0.81	0.5	26.5988	-62
WR	32	1KB	0.151828	0.5	0.5	30.3656	0
WR	64	1KB	0.1508375	0.55	0.5	30.1675	-10
WR	128	1KB	0.1629035	0.67	0.5	32.5807	-34

Conclusion-

From table, we can say that as number of threads increases the latency increases for Random-Write operation and for Random-Read, latency tends to be constant. Disk IOPS tends to be falling

gradually as the number of thread increases for Random-Read operation while IOPS for Random-Write tends to be constant over the number of threads.

Disk Performance Metrics:

Workload = 10GB

Throughput = (Workload / time) / 1000000 Mbps

Theoretical Throughput RS= 372 Mbps Theoretical Throughput WS= 172 Mbps Theoretical Throughput RR= 540 Mbps Theoretical Throughput WR= 410 Mbps

Throughput Efficiency = (Throughput / Theoretical Throughput) * 100

Latency = (execution time * 1000) / 1000000 milliseconds

Theoretical Latency = 0.5 milliseconds

Latency Efficiency = ((Theoretical Latency - Latency) / Theoretical Latency) * 100

Disk IOPS = 1000000 / execution time

Theoretical IOPS = (Theoretical Throughput * 1000000) / 1000

Performance Comparison between Disk Benchmark, IOZONE Benchmark and Theoretical performance

Throughput measured by Disk Benchmark for RR and RS access pattern increases as the number of threads increase which gets double when thread size increases from 1 to 2 and remains approximately similar when threads size increases from 2 to 4. However, for Sequential and Random Write operations, throughput changes. Block size is also a significant factor in determining the throughput. Thread count increase also increases the latency in WR while keeps the constant in RR whereas disk IOPS gets reduced with threads increase in RR while disk IOPS becomes constant in WR. IOZONE efficiency computed for RS and RR is less than the MyDiskBench efficiency but greater for WS and WR operations.

. Network Benchmarking:

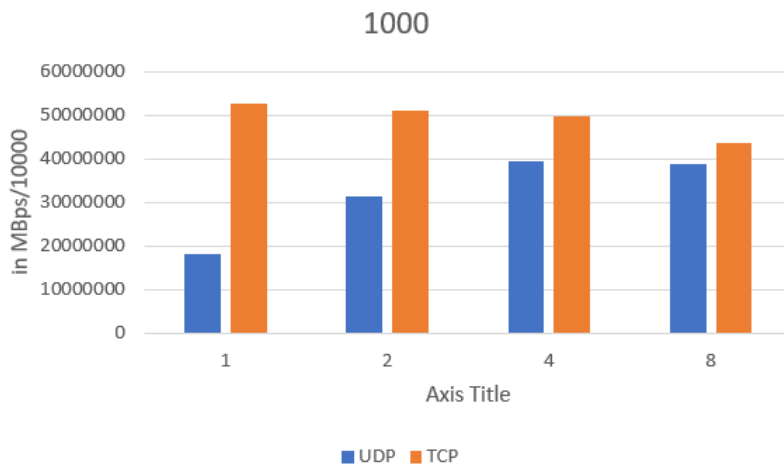
- a. We are using strong scaling to measure N/W speed for loopback network.
- b. Multithreading (up to 8 threads) concept we are using along with socket programming.
- c. Calculating Latency (ms) and Throughput (Mb/sec) for loopback N/W and for this we are using 64 Kb buffer and sending 8k ping pong message .
- d. Both UDP and TCP protocols are used.
- e. Socket programming, we are using in C for both UDP and TCP.

- f. I have not ran this code over Hyperian or any cluster performance matric has been taken through run over my personal system.

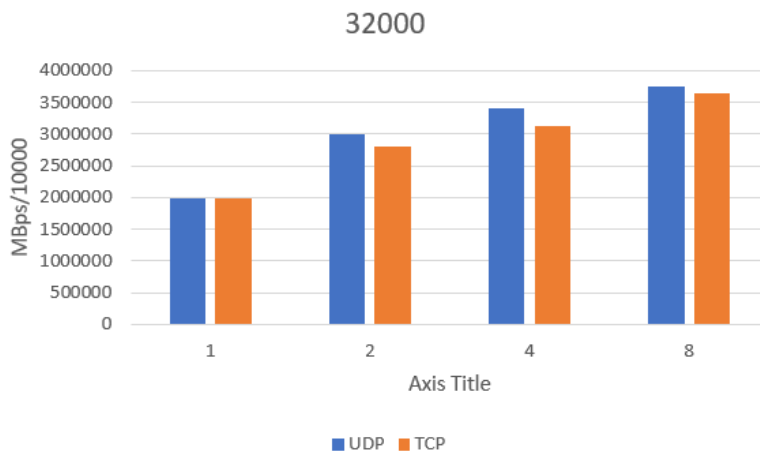
5.1 Performance measurement:

Proto-	Con-	Block	MyNETBench	lperf	Theoretical	MyNETBench	lperf
col	curren- cy	Size	Measured	Measure d	Throughput	Efficiency (%)	Efficienc y
			Throughput	Throughpu t	(Mb/sec)		(%)
			(Mb/sec)	(Mb/sec)			
TCP	1	1KB	1821.8656	2010.41	10000	18.218656	20.1041
TCP	1	32KB	5274.81	2163.52	10000	52.7481	21.6352
TCP	2	1KB	198.715425	3562.63	10000	1.98715425	35.6263
TCP	2	32KB	5098.866	4120.25	10000	50.98866	41.2025
TCP	4	1KB	279.815425	4562.33	10000	2.79815425	45.6233
TCP	4	32KB	4975.81	4765.22	10000	49.7581	47.6522
TCP	8	1KB	312.388925	4412.32	10000	3.12388925	44.1232
TCP	8	32KB	4375.0104	5120.62	10000	43.750104	51.2062
UDP	1	1KB	364.47945	1262.66	10000	3.6447945	12.6266
UDP	1	32KB	198.7197625	6000.35	10000	1.987197625	60.0035
UDP	2	1KB	3148.1254	1432.2	10000	31.481254	14.322
UDP	2	32KB	300.56185	9620.32	10000	3.0056185	96.2032
UDP	4	1KB	3937.0604	3302.32	10000	39.370604	33.0232
UDP	4	32KB	341.34975	8010.32	10000	3.4134975	80.1032
UDP	8	1KB	3874.2412	1595.36	10000	38.742412	15.9536
UDP	8	32KB	374.336425	9162.35	10000	3.74336425	91.6235

TCP/UDP throughput plot : Throughput in Mb/sec vs Concurrency at 1000 B
blocksize:

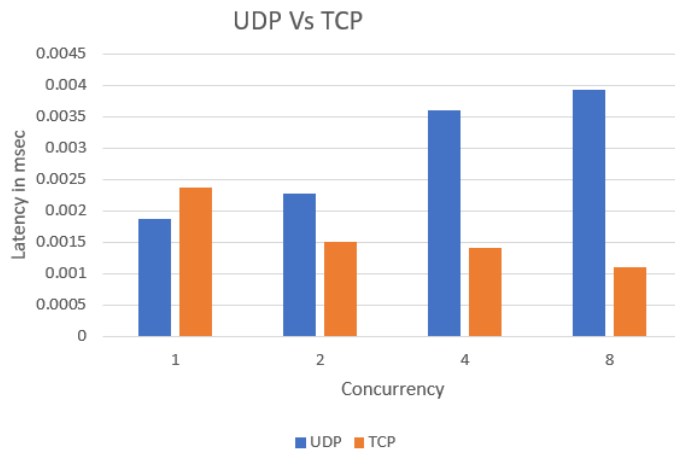


TCP/UPD throughput plot : Throughput in Mb/sec vs Concurrency at 32000 B blocksize:



As from the graph , it can be seen that TCP throughput increase with the increase in thread count but reaches a maximum limit after concurrency level 2. This is due to the sequential copy of data from client buffer to server buffer which involves mutex locks and I have used 1GB buffer at Server and Client side to store the data which is copied in chunks of 1kb or 32kb . Another reason, could be less bandwidth on the hyperion network at the time when the test was run. As UDP follows a connection less service ,there are packet drops involved when sending data.

TCP/UDP latency plot : Latency in us vs block size



It can be seen clearly that for TCP as the no of threads increases , the latency decreases with increase in concurrency. This is not always true for UDP as there is variation in the latency from 1 to 8 threads.

Network Benchmark Metrics:

Workload = 1GB

Throughput = $\frac{((\text{Workload} * 100) / \text{time})}{1048576}$

Theoretical Throughput TCP= 10000 Mbps

Theoretical Throughput UDP= 10000 Mbps

Throughput Efficiency = $\frac{(\text{Throughput} / \text{Theoretical Throughput}) * 100}{100}$

Latency = $\frac{(\text{execution wall time} * 1000)}{1000000}$ milliseconds

Theoretical Latency = 0.0007 milliseconds

Latency Efficiency = $\frac{(\text{Theoretical Latency} - \text{Latency})}{\text{Theoretical Latency}} * 100$

Conclusion:

This performance benchmark on TCP and UDP verifies that , better throughput can be achieved with UDP than TCP due to continuous packet stream between Server and Client with either of them not waiting for any ACK packet from the other, which is the case with TCP. This ACK of sent packets in each TCP opened window delays the transfer of next stream of packets.